

pykalman

Welcome to [pykalman](#), the dead-simple Kalman Filter, Kalman Smoother, and EM library for Python:

```
>>> from pykalman import KalmanFilter
>>> import numpy as np
>>> kf = KalmanFilter(transition_matrices = [[1, 1], [0, 1]], observation_matrices = [[1, 0], [0, 1]], initial_state_mean = 0, initial_state_covariance = 1, initial_state_covariance = 1)
>>> measurements = np.asarray([[1,0], [0,0], [0,1]]) # 3 observations
>>> kf = kf.em(measurements, n_iter=5)
>>> (filtered_state_means, filtered_state_covariances) = kf.filter(measurements)
>>> (smoothed_state_means, smoothed_state_covariances) = kf.smooth(measurements)
```

Also included is support for missing measurements:

```
>>> from numpy import ma
>>> measurements = ma.asarray(measurements)
>>> measurements[1] = ma.masked # measurement at timestep 1 is unknown
>>> kf = kf.em(measurements, n_iter=5)
>>> (filtered_state_means, filtered_state_covariances) = kf.filter(measurements)
>>> (smoothed_state_means, smoothed_state_covariances) = kf.smooth(measurements)
```

And for the non-linear dynamics via the **UnscentedKalmanFilter**:

```
>>> from pykalman import UnscentedKalmanFilter
>>> ukf = UnscentedKalmanFilter(lambda x: x + np.sin(x), lambda x: x, initial_state_mean = 0, initial_state_covariance = 1)
>>> (filtered_state_means, filtered_state_covariances) = ukf.filter([0])
>>> (smoothed_state_means, smoothed_state_covariances) = ukf.smooth([0])
```

Installation

For a quick installation:

```
$ easy_install pykalman
```

pykalman depends on the following modules,

- **numpy** (for core functionality)
- **scipy** (for core functionality)
- **Sphinx** (for generating documentation)
- **numpydoc** (for generating documentation)
- **nose** (for running tests)

All of these and **pykalman** can be installed using **easy_install**:

```
$ easy_install numpy scipy Sphinx numpydoc nose pykalman
```

Alternatively, you can get the latest and greatest from [github](#):

```
$ git clone git@github.com:pykalman/pykalman.git pykalman
$ cd pykalman
$ sudo python setup.py install
```

Kalman Filter User's Guide

The Kalman Filter is a unsupervised algorithm for tracking a single object in a continuous state space. Given a sequence of noisy measurements, the Kalman Filter is able to recover the “true state” of the underling object being tracked. Common uses for the Kalman Filter include radar and sonar tracking and state estimation in robotics.

The advantages of Kalman Filter are:

- No need to provide labeled training data
- Ability to handle noisy observations

The disadvantages are:

- Computational complexity is cubic in the size of the state space
- Parameter optimization is non-convex and can thus only find local optima
- Inability to cope with non-Gaussian noise

Basic Usage

This module implements two algorithms for tracking: the Kalman Filter and Kalman Smoother. In addition, model parameters which are traditionally specified by hand can also be learned by the implemented EM algorithm without any labeled training data. All three algorithms are contained in the `KalmanFilter` class in this module.

In order to apply the Kalman Smoother, one need only specify the size of the state and observation space. This can be done directly by setting `n_dim_state` or `n_dim_obs` or indirectly by specifying an initial value for any of the model parameters from which the former can be derived:

```
>>> from pykalman import KalmanFilter
>>> kf = KalmanFilter(initial_state_mean=0, n_dim_obs=2)
```

The traditional Kalman Filter assumes that model parameters are known beforehand. The `KalmanFilter` class however can learn parameters using `KalmanFilter.em()` (fitting is optional). Then the hidden sequence of states can be predicted using `KalmanFilter.smooth()`:

```
>>> measurements = [[1,0], [0,0], [0,1]]
>>> kf.em(measurements).smooth([[2,0], [2,1], [2,2]])[0]
array([[ 0.85819709],
       [ 1.77811829],
       [ 2.19537816]])
```

The Kalman Filter is parameterized by 3 arrays for state transitions, 3 for measurements, and 2 more for initial conditions. Their names and function are described in the next section.

See also:

`examples/standard/plot_sin.py`
Tracking a sine signal

Choosing Parameters

Unlike most other algorithms, the Kalman Filter and Kalman Smoother are traditionally used with parameters already given. The `KalmanFilter` class can thus be initialized with any subset of the usual model parameters and used without fitting. Sensible default values are given for all unspecified parameters (zeros for all 1-dimensional arrays and identity matrices for all 2-dimensional arrays).

A Kalman Filter/Smoother is fully specified by its initial conditions (`initial_state_mean` and `initial_state_covariance`), its transition parameters (`transition_matrices`, `transition_offsets`, `transition_covariance`), and its observation parameters (`observation_matrices`, `observation_offsets`, `observation_covariance`). These parameters define a probabilistic model from which the unobserved states and observed measurements are assumed to be sampled from. The following code illustrates in one dimension what this process is.

```
from scipy.stats import norm
import numpy as np
states = np.zeros((n_timesteps, n_dim_state))
measurements = np.zeros((n_timesteps, n_dim_obs))
for t in range(n_timesteps-1):
    if t == 0:
        states[t] = norm.rvs(initial_state_mean, np.sqrt(initial_state_
        measurements[t] = (
            np.dot(observation_matrices[t], states[t])
            + observation_offsets[t]
            + norm.rvs(0, np.sqrt(observation_covariance))
        )
    states[t+1] = (
        np.dot(transition_matrices[t], states[t])
        + transition_offsets[t]
        + norm.rvs(0, np.sqrt(transition_covariance))
```

```
)
measurements[t+1] = (
    np.dot(observation_matrices[t+1], states[t+1])
    + observation_offsets[t+1]
    + norm.rvs(np.sqrt(observation_covariance))
)
```

The selection of these variables is not an easy one, and, as shall be explained in the section on fitting, should not be left to `KalmanFilter.em()` alone. If one ignores the random noise, the parameters dictate that *the next state and the current measurement should be an affine function of the current state*. The additive noise term is then simply a way to deal with unaccounted error.

A simple example to illustrate the model parameters is a free falling ball in one dimension. The state vector can be represented by the position, velocity, and acceleration of the ball, and the transition matrix is defined by the equation:

```
position[t+dt] = position[t] + velocity[t] dt + 0.5 acceleration[t] dt2
```

Taking the zeroth, first, and second derivative of the above equation with respect to dt gives the rows of transition matrix:

```
A = np.array([[1, t, 0.5 * (t**2)],
              [0, 1, t],
              [0, 0, 1]])
```

We may also set the transition offset to zero for the position and velocity components and -9.8 for the acceleration component in order to account for gravity's pull.

It is often very difficult to guess what appropriate values are for for the transition and observation covariance, so it is common to use some constant multiplied by the identity matrix. Increasing this constant is equivalent to saying you believe there is more noise in the system. This constant is the amount of variance you expect to see along each dimension during state transitions and measurements, respectively.

Inferring States

The `KalmanFilter` class comes equipped with two algorithms for prediction: the Kalman Filter and the Kalman Smoother. While the former can be updated recursively (making it ideal for online state estimation), the latter can only be done in batch. These two algorithms are accessible via `KalmanFilter.filter()`, `KalmanFilter.filter_update()`, and `KalmanFilter.smooth()`.

Functionally, Kalman Smoother should always be preferred. Unlike the Kalman Filter, the Smoother is able to incorporate “future” measurements as well as past ones at the same computational cost of $O(Td^3)$ where T is the number of time steps and d is the

dimensionality of the state space. The only reason to prefer the Kalman Filter over the Smoother is in its ability to incorporate new measurements in an online manner:

```
>>> means, covariances = kf.filter(measurements)
>>> next_mean, next_covariance = kf.filter_update(
    means[-1], covariances[-1], new_measurement
)
```

Both the Kalman Filter and Kalman Smoother are able to use parameters which vary with time. In order to use this, one need only pass in an array `n_timesteps` in length along its first axis:

```
>>> transition_offsets = [[-1], [0], [1], [2]]
>>> kf = KalmanFilter(transition_offsets=transition_offsets, n_dim_obs=4)
```

See also:

`examples/standard/plot_online.py`

Online State Estimation

`examples/standard/plot_filter.py`

Filtering and Smoothing

Optimizing Parameters

In addition to the Kalman Filter and Kalman Smoother, the `KalmanFilter` class implements the Expectation-Maximization algorithm. This iterative algorithm is a way to maximize the likelihood of the observed measurements (recall the probabilistic model induced by the model parameters), which is unfortunately a non-convex optimization problem. This means that even when the EM algorithm converges, there is no guarantee that it has converged to an optimal value. Thus it is important to select good initial parameter values.

A second consideration when using the EM algorithm is that the algorithm lacks regularization, meaning that parameter values may diverge to infinity in order to make the measurements more likely. Thus it is important to choose *which* parameters to optimize via the `em_vars` parameter of `KalmanFilter`. For example, in order to only optimize the transition and observation covariance matrices, one may instantiate `KalmanFilter` like so:

```
>>> kf = KalmanFilter(em_vars=['transition_covariance', 'observation_covariance'],
```

It is customary to optimize only the `transition_covariance`, `observation_covariance`, `initial_state_mean`, and `initial_state_covariance`,

which is the default when `em_vars` is unspecified. In order to avoid overfitting, it is also possible to specify the number of iterations of the EM algorithm to run during fitting:

```
>>> kf.em(X, n_iter=5)
```

Each iteration of the EM algorithm requires running the Kalman Smoother anew, so its computational complexity is $O(Tnd^3)$ where T is the number of time steps, n is the number of iterations, and d is the size of the state space.

See also:

`examples/standard/plot_em.py`
Using the EM Algorithm

Missing Measurements

In real world systems, it is common to have sensors occasionally fail. The Kalman Filter, Kalman Smoother, and EM algorithm are all equipped to handle this scenario. To make use of it, one only need apply a NumPy mask to the measurement at the missing time step:

```
>>> from numpy import ma
>>> X = ma.array([1,2,3])
>>> X[1] = ma.masked # hide measurement at time step 1
>>> kf.em(X).smooth(X)
```

See also:

`examples/standard/plot_missing.py`
State Estimation with Missing Observations

Mathematical Formulation

In order to understand when the algorithms in this module will be effective, it is important to understand what assumptions are being made. To make notation concise, we refer to the hidden states as \mathbf{x}_t , the measurements as \mathbf{z}_t , and the parameters of the `KalmanFilter` class as follows,

Parameter Name	Notation
<code>initial_state_mean</code>	μ_0
<code>initial_state_covariance</code>	Σ_0
<code>transition_matrices</code>	A
<code>transition_offsets</code>	b

<i>transition_covariance</i>	Q
<i>observation_matrices</i>	C
<i>observation_offsets</i>	b
<i>observation_covariance</i>	R

In words, the Linear-Gaussian model assumes that for all time steps $t = 0, \dots, T - 1$ (here, T is the number of time steps),

- x_0 is distributed according to a Gaussian distribution
- x_{t+1} is an affine transformation of x_t and additive Gaussian noise
- z_t is an affine transformation of x_t and additive Gaussian noise

These assumptions imply that that x_t is always a Gaussian distribution, even when z_t is observed. If this is the case, the distribution of $x_t|z_{1:t}$ and $x_t|z_{1:T-1}$ are completely specified by the parameters of the Gaussian distribution, namely its *mean* and *covariance*. The Kalman Filter and Kalman Smoother calculate these values, respectively.

Formally, the Linear-Gaussian Model assumes that states and measurements are generated in the following way,

$$\begin{aligned}
 x_0 &\sim \text{Gaussian}(\mu_0, \Sigma_0) \\
 x_{t+1} &= A_t x_t + b_t + \epsilon_{t+1}^1 \\
 y_t &= C_t x_t + d_t + \epsilon_t^2 \\
 \epsilon_t^1 &\sim \text{Gaussian}(0, Q) \\
 \epsilon_t^2 &\sim \text{Gaussian}(0, R)
 \end{aligned}$$

The Gaussian distribution is characterized by its single mode and exponentially decreasing tails, meaning that the Kalman Filter and Kalman Smoother work best if one is able to guess fairly well the vicinity of the next state given the present, but cannot say *exactly* where it will be. On the other hand, these methods will fail if there are multiple, disconnected areas where the next state could be, such as if a car turns one of three ways at an intersection.

References:

- Abbeel, Pieter. "Maximum Likelihood, EM". <http://www.cs.berkeley.edu/~pabbeel/cs287-fa11/>
- Yu, Byron M. and Shenoy, Krishna V. and Sahani, Maneesh. "Derivation of Kalman Filtering and Smoothing Equations". http://www.ece.cmu.edu/~byronyu/papers/derive_ks.pdf
- Ghahramani, Zoubin and Hinton, Geoffrey E. "Parameter Estimation for Linear Dynamical Systems." <http://mlg.eng.cam.ac.uk/zoubin/course04/tr-96-2.pdf>
- Welling, Max. "The Kalman Filter". http://www.cs.toronto.edu/~welling/classnotes/papers_class/KF.ps.gz

Unscented Kalman Filter User's Guide

Like the Kalman Filter, the Unscented Kalman Filter is an unsupervised algorithm for tracking a single target in a continuous state space. The difference is that while the Kalman Filter restricts dynamics to affine functions, the Unscented Kalman Filter is designed to operate under arbitrary dynamics.

The advantages of the Unscented Kalman Filter implemented here are:

- Ability to handle non-affine state transition and observation functions
- Ability to handle not-quite-Gaussian noise models
- Same computational complexity as the standard Kalman Filter

The disadvantages are:

- No method for learning parameters
- Lack of theoretical guarantees on performance
- Inability to handle extremely non-Gaussian noise

Basic Usage

Like `KalmanFilter`, two methods are provided in `UnscentedKalmanFilter` for tracking targets: `UnscentedKalmanFilter.filter()` and `UnscentedKalmanFilter.smooth()`. At this point no algorithms have been implemented for inferring parameters, so they must be specified by hand at instantiation.

In order to apply these algorithms, one must specify a subset of the following,

Variable Name	Mathematical Notation	Default
<code>transition_functions</code>	f_t	state plus noise
<code>observation_functions</code>	g_t	state plus noise
<code>transition_covariance</code>	Q	identity
<code>observation_covariance</code>	R	identity
<code>initial_state_mean</code>	μ_0	zero
<code>initial_state_covariance</code>	Σ_0	identity

If parameters are left unspecified, they will be replaced by their defaults. One also has the option of simply specifying `n_dim_state` or `n_dim_obs` if the size of the state or observation space cannot be inferred directly.

The state transition function and observation function have replaced the transition matrix/offset and observation matrix/offset from the original `KalmanFilter`, respectively. Both must take in the current state and some Gaussian-sampled noise

and return the next state/current observation. For example, if noise were multiplicative instead of additive, the following would be valid:

```
>>> def f(current_state, transition_noise):
...     return current_state * transition_noise
...
>>> def g(current_state, observation_noise):
...     return current_state * observation_noise
```

Once defined, the `UnscentedKalmanFilter` can be used to extract estimated state and covariance matrices over the hidden state:

```
>>> from pykalman import UnscentedKalmanFilter
>>> import numpy as np
>>> def f(state, noise):
...     return state + np.sin(noise)
...
>>> def g(state, noise):
...     return state + np.cos(noise)
...
>>> ukf = UnscentedKalmanFilter(f, g, R=0.1)
>>> ukf.smooth([0, 1, 2])[0]
array([[ -0.94034641],
       [  0.05002316],
       [  1.04502498]])
```

If the `UnscentedKalmanFilter` is instantiated with an array of functions for `transition_functions` or `observation_functions`, then the function is assumed to vary with time. Currently there is no support for time-varying covariance matrices.

Which Unscented Kalman Filter is for Me?

Though only `UnscentedKalmanFilter` was mentioned in the previous section, there exists another class specifically designed for the case when noise is additive, `AdditiveUnscentedKalmanFilter`. While more restrictive, this class offers reduced computational complexity ($O(Tn^3)$ vs. $O(T(2n+m)^3)$ for state space with dimensionality n , observation space with dimensionality m) and better numerical stability. When at all possible, the `AdditiveUnscentedKalmanFilter` should be preferred to its counterpart.

To reflect the restriction on how noise is integrated, the `AdditiveUnscentedKalmanFilter` uses state transition and observation functions with slightly different arguments:

```
def f(current_state):
    ...
```

```
def g(current_state):  
    ...
```

Notice that the transition/observation noise is no longer an argument. Its effect will be taken care of at later points in the algorithm without any need for your explicit input.

Finally, users should note that the `UnscentedKalmanFilter` can potentially suffer from collapse of the covariance matrix to zero. Algorithmically, this means that the `UnscentedKalmanFilter` is one hundred percent sure of the state and that no noise is left in the system. In order to avoid this, one must ensure that even for small amounts of noise, `transition_functions` and `observation_functions` output different values for the same current state.

Choosing Parameters

The majority of advice on choosing parameters in Kalman Filter section apply to the Unscented Kalman Filter except that there is no method for learning parameters and the following code snippet defines the probabilistic model the Unscented Kalman Filter (approximately) solves,

```
from scipy.stats import norm  
import numpy as np  
states = np.zeros((n_timesteps, n_dim_state))  
measurements = np.zeros((n_timesteps, n_dim_obs))  
for t in range(n_timesteps-1):  
    if t == 0:  
        states[t] = norm.rvs(initial_state_mean, np.sqrt(initial_state_  
        measurements[t] = (  
            observation_function(  
                states[t],  
                norm.rvs(0, np.sqrt(observation_covariance))  
            )  
        )  
    states[t+1] = (  
        transition_function(  
            states[t],  
            norm.rvs(0, np.sqrt(transition_covariance))  
        )  
    )  
    measurements[t+1] = (  
        observation_function(  
            states[t+1],  
            norm.rvs(0, np.sqrt(observation_covariance))  
        )  
    )
```

Missing Measurements

The `UnscentedKalmanFilter` and `AdditiveUnscentedKalmanFilter` have the same support for missing measurements that the original `KalmanFilter` class supports. Usage is precisely the same.

Class Reference

KalmanFilter

```
class pykalman.KalmanFilter(transition_matrices=None,
                             observation_matrices=None, transition_covariance=None,
                             observation_covariance=None, transition_offsets=None, observation_offsets=None,
                             initial_state_mean=None, initial_state_covariance=None, random_state=None,
                             em_vars=['transition_covariance', 'observation_covariance', 'initial_state_mean',
                                      'initial_state_covariance'], n_dim_state=None, n_dim_obs=None)
```

Implements the Kalman Filter, Kalman Smoother, and EM algorithm.

This class implements the Kalman Filter, Kalman Smoother, and EM Algorithm for a Linear Gaussian model specified by,

$$\begin{aligned}x_{t+1} &= A_t x_t + b_t + \text{Normal}(0, Q_t) \\ z_t &= C_t x_t + d_t + \text{Normal}(0, R_t)\end{aligned}$$

The Kalman Filter is an algorithm designed to estimate $P(x_t | z_{0:t})$. As all state transitions and observations are linear with Gaussian distributed noise, these distributions can be represented exactly as Gaussian distributions with mean *filtered_state_means[t]* and covariances *filtered_state_covariances[t]*.

Similarly, the Kalman Smoother is an algorithm designed to estimate $P(x_t | z_{0:T-1})$.

The EM algorithm aims to find for $\theta = (A, b, C, d, Q, R, \mu_0, \Sigma_0)$

$$\max_{\theta} P(z_{0:T-1}; \theta)$$

If we define $L(x_{0:T-1}, \theta) = \log P(z_{0:T-1}, x_{0:T-1}; \theta)$, then the EM algorithm works by iteratively finding,

$$P(x_{0:T-1} | z_{0:T-1}, \theta_i)$$

then by maximizing,

$$\theta_{i+1} = \arg \max_{\theta} \mathbb{E}_{x_{0:T-1}} [L(x_{0:T-1}, \theta) | z_{0:T-1}, \theta_i]$$

Parameters : **transition_matrices** : [n_timesteps-1, n_dim_state, n_dim_state] or [n_dim_state, n_dim_state] array-like

Also known as A . state transition matrix between times t and $t+1$ for t in $[0...n_timesteps-2]$

observation_matrices : $[n_timesteps, n_dim_obs, n_dim_obs]$ or $[n_dim_obs, n_dim_obs]$ array-like

Also known as C . observation matrix for times $[0...n_timesteps-1]$

transition_covariance : $[n_dim_state, n_dim_state]$ array-like

Also known as Q . state transition covariance matrix for times $[0...n_timesteps-2]$

observation_covariance : $[n_dim_obs, n_dim_obs]$ array-like

Also known as R . observation covariance matrix for times $[0...n_timesteps-1]$

transition_offsets : $[n_timesteps-1, n_dim_state]$ or $[n_dim_state]$ array-like

Also known as b . state offsets for times $[0...n_timesteps-2]$

observation_offsets : $[n_timesteps, n_dim_obs]$ or $[n_dim_obs]$ array-like

Also known as d . observation offset for times $[0...n_timesteps-1]$

initial_state_mean : $[n_dim_state]$ array-like

Also known as μ_0 . mean of initial state distribution

initial_state_covariance : $[n_dim_state, n_dim_state]$ array-like

Also known as Σ_0 . covariance of initial state distribution

random_state : optional, numpy random state

random number generator used in sampling

em_vars : optional, subset of $['transition_matrices', 'observation_matrices', 'transition_offsets', 'observation_offsets', 'transition_covariance', 'observation_covariance', 'initial_state_mean', 'initial_state_covariance']$ or 'all'

if *em_vars* is an iterable of strings only variables in *em_vars* will be estimated using EM. if *em_vars* == 'all', then all variables will be estimated.

n_dim_state: optional, integer :

the dimensionality of the state space. Only meaningful when you do not specify initial values for *transition_matrices*, *transition_offsets*, *transition_covariance*, *initial_state_mean*, or *initial_state_covariance*.

n_dim_obs: optional, integer :

the dimensionality of the observation space. Only meaningful when you do not specify initial values for *observation_matrices*, *observation_offsets*, or *observation_covariance*.

Methods

em (X [, y, n_iter, em_vars])	Apply the EM algorithm
filter (X)	Apply the Kalman Filter
filter_update (filtered_state_mean, ... [, ...])	Update a Kalman Filter state estimate
loglikelihood (X)	Calculate the log likelihood of all observations
sample (n_timesteps [, initial_state, ...])	Sample a state sequence <i>n_timesteps</i> timesteps in length.
smooth (X)	Apply the Kalman Smoother

em(X, y=None, n_iter=10, em_vars=None)

Apply the EM algorithm

Apply the EM algorithm to estimate all parameters specified by *em_vars*. Note that all variables estimated are assumed to be constant for all time. See *_em()* for details.

Parameters : **X** : [n_timesteps, n_dim_obs] array-like

observations corresponding to times [0...n_timesteps-1]. If X is a masked array and any of *X[t]*'s components is masked, then *X[t]* will be treated as a missing observation.

n_iter : int, optional

number of EM iterations to perform

em_vars : iterable of strings or 'all'

variables to perform EM over. Any variable not appearing here is left untouched.

filter(X)

Apply the Kalman Filter

Apply the Kalman Filter to estimate the hidden state at time t for $t = [0 \dots n_{\text{timesteps}} - 1]$ given observations up to and including time t . Observations are assumed to correspond to times $[0 \dots n_{\text{timesteps}} - 1]$. The output of this method corresponding to time $n_{\text{timesteps}} - 1$ can be used in `KalmanFilter.filter_update()` for online updating.

Parameters : **X** : [n_timesteps, n_dim_obs] array-like

observations corresponding to times $[0 \dots n_{\text{timesteps}} - 1]$. If X is a masked array and any of $X[t]$ is masked, then $X[t]$ will be treated as a missing observation.

Returns : **filtered_state_means** : [n_timesteps, n_dim_state]

mean of hidden state distributions for times $[0 \dots n_{\text{timesteps}} - 1]$ given observations up to and including the current time step

filtered_state_covariances : [n_timesteps, n_dim_state, n_dim_state] array

covariance matrix of hidden state distributions for times $[0 \dots n_{\text{timesteps}} - 1]$ given observations up to and including the current time step

filter_update(*filtered_state_mean*, *filtered_state_covariance*,
observation=None, *transition_matrix=None*, *transition_offset=None*,
transition_covariance=None, *observation_matrix=None*,
observation_offset=None, *observation_covariance=None*)

Update a Kalman Filter state estimate

Perform a one-step update to estimate the state at time $t + 1$ give an observation at time $t + 1$ and the previous estimate for time t given

observations from times $[0...t]$. This method is useful if one wants to track an object with streaming observations.

Parameters : **filtered_state_mean** : $[n_dim_state]$ array

mean estimate for state at time t given
observations from times $[1...t]$

filtered_state_covariance : $[n_dim_state, n_dim_state]$
array

covariance of estimate for state at time t given
observations from times $[1...t]$

observation : $[n_dim_obs]$ array or None

observation from time $t+1$. If *observation* is a
masked array and any of *observation*'s
components are masked or if *observation* is
None, then *observation* will be treated as a
missing observation.

transition_matrix : optional, $[n_dim_state, n_dim_state]$
array

state transition matrix from time t to $t+1$. If
unspecified, *self.transition_matrices* will be used.

transition_offset : optional, $[n_dim_state]$ array

state offset for transition from time t to $t+1$. If
unspecified, *self.transition_offset* will be used.

transition_covariance : optional, $[n_dim_state,$
 $n_dim_state]$ array

state transition covariance from time t to $t+1$. If
unspecified, *self.transition_covariance* will be
used.

observation_matrix : optional, $[n_dim_obs, n_dim_state]$
array

observation matrix at time $t+1$. If unspecified,
self.observation_matrices will be used.

observation_offset : optional, $[n_dim_obs]$ array

observation offset at time t+1. If unspecified, *self.observation_offset* will be used.

observation_covariance : optional, [n_dim_obs, n_dim_obs] array

observation covariance at time t+1. If unspecified, *self.observation_covariance* will be used.

Returns : **next_filtered_state_mean** : [n_dim_state] array

mean estimate for state at time t+1 given observations from times [1...t+1]

next_filtered_state_covariance : [n_dim_state, n_dim_state] array

covariance of estimate for state at time t+1 given observations from times [1...t+1]

loglikelihood(X)

Calculate the log likelihood of all observations

Parameters : **X** : [n_timesteps, n_dim_obs] array

observations for time steps [0...n_timesteps-1]

Returns : **likelihood** : float

likelihood of all observations

sample(n_timesteps, initial_state=None, random_state=None)

Sample a state sequence $n_{\text{timesteps}}$ timesteps in length.

Parameters : **n_timesteps** : int

number of timesteps

Returns : **states** : [n_timesteps, n_dim_state] array

hidden states corresponding to times [0...n_timesteps-1]

observations : [n_timesteps, n_dim_obs] array

observations corresponding to times
[0...n_timesteps-1]

smooth(X)

Apply the Kalman Smoother

Apply the Kalman Smoother to estimate the hidden state at time t for $t = [0...n_{\text{timesteps}} - 1]$ given all observations. See `_smooth()` for more complex output

Parameters : **X** : [n_timesteps, n_dim_obs] array-like

observations corresponding to times
[0...n_timesteps-1]. If X is a masked array and
any of $X[t]$ is masked, then $X[t]$ will be treated as
a missing observation.

Returns : **smoothed_state_means** : [n_timesteps, n_dim_state]

mean of hidden state distributions for times
[0...n_timesteps-1] given all observations

smoothed_state_covariances : [n_timesteps,
n_dim_state]

covariances of hidden state distributions for
times [0...n_timesteps-1] given all observations

UnscentedKalmanFilter

```
class pykalman.UnscentedKalmanFilter(transition_functions=None,  
observation_functions=None, transition_covariance=None,  
observation_covariance=None, initial_state_mean=None,  
initial_state_covariance=None, n_dim_state=None, n_dim_obs=None,  
random_state=None)
```

Implements the General (aka Augmented) Unscented Kalman Filter governed by the following equations,

$$\begin{aligned}x_0 &\sim \text{Normal}(\mu_0, \Sigma_0) \\x_{t+1} &= f_t(x_t, \text{Normal}(0, Q)) \\z_t &= g_t(x_t, \text{Normal}(0, R))\end{aligned}$$

Notice that although the input noise to the state transition equation and the observation equation are both normally distributed, any non-linear transformation may be applied afterwards. This allows for greater generality, but at the expense of computational complexity. The complexity of `UnscentedKalmanFilter.filter`

$()$ is $O(T(2n + m)^3)$ where T is the number of time steps, n is the size of the state space, and m is the size of the observation space.

If your noise is simply additive, consider using the [AdditiveUnscentedKalmanFilter](#)

Parameters : **transition_functions** : function or [n_timesteps-1] array of functions

transition_functions[t] is a function of the state and the transition noise at time t and produces the state at time t+1. Also known as f_t .

observation_functions : function or [n_timesteps] array of functions

observation_functions[t] is a function of the state and the observation noise at time t and produces the observation at time t. Also known as g_t .

transition_covariance : [n_dim_state, n_dim_state] array

transition noise covariance matrix. Also known as Q .

observation_covariance : [n_dim_obs, n_dim_obs] array

observation noise covariance matrix. Also known as R .

initial_state_mean : [n_dim_state] array

mean of initial state distribution. Also known as μ_0

initial_state_covariance : [n_dim_state, n_dim_state] array

covariance of initial state distribution. Also known as Σ_0

n_dim_state: optional, integer :

the dimensionality of the state space. Only meaningful when you do not specify initial values for *transition_covariance*, or *initial_state_mean*, *initial_state_covariance*.

n_dim_obs: optional, integer :

the dimensionality of the observation space. Only meaningful when you do not specify initial values for *observation_covariance*.

random_state : optional, int or RandomState

seed for random sample generation

Methods

filter (Z)	Run Unscented Kalman Filter
sample (n_timesteps [, initial_state, ...])	Sample from model defined by the Unscented Kalman Filter
smooth (Z)	Run Unscented Kalman Smoother

filter(Z)

Run Unscented Kalman Filter

Parameters : **Z** : [n_timesteps, n_dim_state] array

Z[t] = observation at time t. If Z is a masked array and any of Z[t]'s elements are masked, the observation is assumed missing and ignored.

Returns : **filtered_state_means** : [n_timesteps, n_dim_state] array

filtered_state_means[t] = mean of state distribution at time t given observations from times [0, t]

filtered_state_covariances : [n_timesteps, n_dim_state, n_dim_state] array

filtered_state_covariances[t] = covariance of state distribution at time t given observations from times [0, t]

sample(n_timesteps, initial_state=None, random_state=None)

Sample from model defined by the Unscented Kalman Filter

Parameters : **n_timesteps** : int

number of time steps

initial_state : optional, [n_dim_state] array

initial state. If unspecified, will be sampled from initial state distribution.

random_state : optional, int or Random

random number generator

smooth(Z)

Run Unscented Kalman Smoother

Parameters : **Z** : [n_timesteps, n_dim_state] array

$Z[t]$ = observation at time t . If Z is a masked array and any of $Z[t]$'s elements are masked, the observation is assumed missing and ignored.

Returns : **smoothed_state_means** : [n_timesteps, n_dim_state] array

$\text{filtered_state_means}[t]$ = mean of state distribution at time t given observations from times $[0, n_timesteps-1]$

smoothed_state_covariances : [n_timesteps, n_dim_state, n_dim_state] array

$\text{filtered_state_covariances}[t]$ = covariance of state distribution at time t given observations from times $[0, n_timesteps-1]$

AdditiveUnscentedKalmanFilter

class pykalman.**AdditiveUnscentedKalmanFilter**

(*transition_functions=None, observation_functions=None, transition_covariance=None, observation_covariance=None, initial_state_mean=None, initial_state_covariance=None, n_dim_state=None, n_dim_obs=None, random_state=None*)

Implements the Unscented Kalman Filter with additive noise. Observations are assumed to be generated from the following process,

$$\begin{aligned}x_0 &\sim \text{Normal}(\mu_0, \Sigma_0) \\x_{t+1} &= f_t(x_t) + \text{Normal}(0, Q) \\z_t &= g_t(x_t) + \text{Normal}(0, R)\end{aligned}$$

While less general the general-noise Unscented Kalman Filter, the Additive version is more computationally efficient with complexity $O(Tn^3)$ where T is the number of time steps and n is the size of the state space.

Parameters :

transition_functions : function or [n_timesteps-1] array of functions

transition_functions[t] is a function of the state at time t and produces the state at time t+1. Also known as f_t .

observation_functions : function or [n_timesteps] array of functions

observation_functions[t] is a function of the state at time t and produces the observation at time t. Also known as g_t .

transition_covariance : [n_dim_state, n_dim_state] array

transition noise covariance matrix. Also known as Q .

observation_covariance : [n_dim_obs, n_dim_obs] array

observation noise covariance matrix. Also known as R .

initial_state_mean : [n_dim_state] array

mean of initial state distribution. Also known as μ_0 .

initial_state_covariance : [n_dim_state, n_dim_state] array

covariance of initial state distribution. Also known as Σ_0 .

n_dim_state: optional, integer :

the dimensionality of the state space. Only meaningful when you do not specify initial values for *transition_covariance*, or *initial_state_mean*, *initial_state_covariance*.

n_dim_obs: optional, integer :

the dimensionality of the observation space. Only meaningful when you do not specify initial values for *observation_covariance*.

random_state : optional, int or RandomState

seed for random sample generation

Methods

<code>filter(Z)</code>	Run Unscented Kalman Filter
<code>sample</code> (<code>n_timesteps</code> [, <code>initial_state</code> , ...])	Sample from model defined by the Unscented Kalman Filter
<code>smooth(Z)</code>	Run Unscented Kalman Smoother

`filter(Z)`

Run Unscented Kalman Filter

Parameters : **Z** : [n_timesteps, n_dim_state] array

$Z[t]$ = observation at time t . If Z is a masked array and any of $Z[t]$'s elements are masked, the observation is assumed missing and ignored.

Returns : **filtered_state_means** : [n_timesteps, n_dim_state] array

`filtered_state_means[t]` = mean of state distribution at time t given observations from times $[0, t]$

filtered_state_covariances : [n_timesteps, n_dim_state, n_dim_state] array

`filtered_state_covariances[t]` = covariance of state distribution at time t given observations from times $[0, t]$

`sample(n_timesteps, initial_state=None, random_state=None)`

Sample from model defined by the Unscented Kalman Filter

Parameters : **n_timesteps** : int

number of time steps

initial_state : optional, [n_dim_state] array

initial state. If unspecified, will be sampled from initial state distribution.

`smooth(Z)`

Run Unscented Kalman Smoother

Parameters : **Z** : [n_timesteps, n_dim_state] array

$Z[t]$ = observation at time t . If Z is a masked array and any of $Z[t]$'s elements are masked, the observation is assumed missing and ignored.

Returns : **smoothed_state_means** : $[n_timesteps, \quad n_dim_state]$
array

filtered_state_means[t] = mean of state
distribution at time t given observations from
times $[0, n_timesteps-1]$

smoothed_state_covariances : $[n_timesteps,$
 $n_dim_state, n_dim_state]$ array

filtered_state_covariances[t] = covariance of
state distribution at time t given observations
from times $[0, n_timesteps-1]$